

CSCI446/946 Big Data Analytics

Week 5 Advanced Analytical Theory and Methods: Association Rules

School of Computing and Information Technology
University of Wollongong Australia

Advanced Analytical Theory and Methods: **Association Rules**

- Overview of Association Rules
- **Apriori Algorithm**
- Evaluation of Candidate Rules
- An example of rule generation in R
- Validation and Testing
- Diagnostics

All the figures, tables and codes are from the book “[Data Science and Big Data Analytics: Discovering, Analyzing, Visualizing and Presenting Data](#)” unless indicated otherwise.

Advanced Analytical Theory and Methods: Association Rules

- An **unsupervised** learning method
- **Descriptive**, not predictive
- Discover **interesting, hidden** relationship
 - Represented as **rules** or **frequent itemsets**
- Commonly used for **mining** transactions in databases

All the figures, tables and codes are from the book “[Data Science and Big Data Analytics: Discovering, Analyzing, Visualizing and Presenting Data](#)” unless indicated otherwise.

Advanced Analytical Theory and Methods: Association Rules

- It can usually answer the questions like
 - Which products tend to be purchased together?
 - Of those customers who are similar to this person, what products do they tend to buy?
 - Of those customers who have purchased this product, what other products do they tend to view or purchase?

Overview of Association Rules

- Each **transaction** consists of one or more **items**.
- What items are **frequently** purchased together?
- Goal: discover “**interesting**” relationships among the items.



Overview of Association Rules

- Uncovered **rule** is in the form $X \rightarrow Y$
 - meaning that when item X is **observed**, item Y is also **observed**
 - X: left-hand side (**lhs**); Y: right-hand side (**rhs**)
 - What does “**Cereal \rightarrow Milk (90%)**” mean?
 - When cereal is purchased, **90% of the time** milk is also purchased.



Overview of Association Rules

- Also known as “market basket analysis”
 - Each transaction – shopping basket
- Itemset
 - A collection of items or individual entities that contain some kind of relationship
- k-itemset
 - An itemset containing k items
 - {item1, item2, ..., item k }



Overview of Association Rules

- Exhaustively check all possible itemsets?
 - No! The size is exponentially large...
- Apriori algorithm
 - One of the earliest and the most fundamental algorithms for generating association rules.
- Key concept: support
 - For pruning itemsets and controlling the exponential growth of candidate itemsets.

Overview of Association Rules


- Support
 - Given an item X , the support of X is the percentage of transactions that contain X
 - Denoted by $\text{support}(X)$
- Frequent itemset
 - Contains items that appear together often enough
 - Formally, its support \geq a minimum support

Overview of Association Rules

- **Apriori property** (downward closure property)
 - If an itemset is **frequent**, then any **subset** of this itemset must also be **frequent**
 - It provides the **basis** for the Apriori algorithm
- An example: If $\text{support}(\{\text{bread}, \text{jam}\}) = 0.6 \rightarrow$
 $\text{Support}(\{\text{bread}\}) \geq 0.6$ and $\text{Support}(\{\text{jam}\}) \geq 0.6$
- Therefore, if X is infrequent then all supersets that contain X must also be infrequent.

Overview of Association Rules

- Apriori property (downward closure property)



Itemset {A, B, C, D} and its subsets

Apriori Algorithm

- It takes a **bottom-up** iterative approach to uncovering frequent itemsets
 - First, identify all **frequent** items (or **1-itemsets**)
 - The identified frequent 1-itemsets are paired into 2-itemsets to identify **frequent 2-itemsets**
 - **Grow** the size of identified frequent itemsets and **identify** again
 - **Repeat** this process **until** 1) it runs out of support or 2) the itemsets reach a predefined length

Apriori Algorithm

Input

- A transaction database D
- A minimum support threshold δ
- An optional parameter N indicating the maximum length an itemset could reach

```
1  Apriori ( $D, \delta, N$ )
2     $k \leftarrow 1$ 
3     $L_k \leftarrow \{ \text{1-itemsets that satisfy minimum support } \delta \}$ 
4    while  $L_k \neq \emptyset$ 
5      if  $\nexists N \vee (\exists N \wedge k < N)$ 
6         $C_{k+1} \leftarrow \text{candidate itemsets generated from } L_k$ 
7        for each transaction  $t$  in database  $D$  do
8          increment the counts of  $C_{k+1}$  contained in  $t$ 
9         $L_{k+1} \leftarrow \text{candidates in } C_{k+1} \text{ that satisfy minimum support } \delta$ 
10        $k \leftarrow k + 1$ 
11    return  $\bigcup_k L_k$ 
```

Apriori Algorithm

- **Output** of the Apriori algorithm
 - The collection of all the frequent k-itemsets
- A collection of **candidate rules** is formed based on the frequent itemsets uncovered
 - {milk, eggs} may suggest candidate rules
 - {milk} \rightarrow {eggs} and {eggs} \rightarrow {milk}
- Implemented by **apriori()** function in R

[illegible]

Evaluation of Candidate Rules

- How to **evaluate** the **appropriateness** of these candidate rules?
 - Many measures!
 - Confidence, lift, leverage are among the most common.
- **Confidence**
 - The measure of **certainty or trustworthiness** associated with each rule

$$\text{Confidence}(X \rightarrow Y) = \frac{\text{Support}(X \wedge Y)}{\text{Support}(X)}$$

Evaluation of Candidate Rules

- Minimum Confidence

- A predefined threshold to indicate a relationship is “interesting”.
- A higher confidence **could** indicates that the rule $(X \rightarrow Y)$ is more interesting (**be careful...**).
- All the rules can be **ranked** based on **support** or **confidence**

$$Confidence(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X)}$$

Evaluation of Candidate Rules

- Issue with “Confidence”
 - In what cases will we obtain a **high** confidence?
 - It **cannot** tell
 - if a rule contains true implication of the relationship
 - If the rule is purely coincidental

$$Confidence(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X)}$$

Evaluation of Candidate Rules

- Lift

- Measures how many times **more often** X and Y occur together **than expected** if they are statistically **independent** of each other.
- Measures how X and Y are really related rather than **coincidentally** happening together:

$$Lift(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X) * Support(Y)}$$

Evaluation of Candidate Rules

- Lift

- Lift is 1 if X and Y are statistically independent of each other.
- A lift of $X \rightarrow Y$ greater than 1 indicates some usefulness of the rule.
- A larger lift suggests a greater strength of the association between X and Y.

$$Lift(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X) * Support(Y)}$$

Evaluation of Candidate Rules

- Leverage (Pitetsky-Shapiro's)
 - Measures the **difference** in the probability of X and Y appearing together compared to what would be expected if X and Y were **statistically independent** of each other

$$Lift(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X) * Support(Y)}$$

$$Leverage(X \rightarrow Y) = Support(X \wedge Y) - Support(X) * Support(Y)$$

Evaluation of Candidate Rules

- Leverage

- Its value will be zero when X and Y are statistically independent of each other.
- If X and Y have some kind of relationship, the leverage would be greater than zero.

$$Lift(X \rightarrow Y) = \frac{Support(X \wedge Y)}{Support(X) * Support(Y)}$$

$$Leverage(X \rightarrow Y) = Support(X \wedge Y) - Support(X) * Support(Y)$$

Evaluation of Candidate Rules

- Four measures

- Support, Confidence, Lift, and Leverage
- A **high-confidence** rule can sometimes be **misleading**.
- Lift and leverage not only ensure interesting rules but also filter out **coincidental** rules.

$$\text{Confidence}(X \rightarrow Y) = \frac{\text{Support}(X \wedge Y)}{\text{Support}(X)} \quad \text{Lift}(X \rightarrow Y) = \frac{\text{Support}(X \wedge Y)}{\text{Support}(X) * \text{Support}(Y)}$$

$$\text{Leverage}(X \rightarrow Y) = \text{Support}(X \wedge Y) - \text{Support}(X) * \text{Support}(Y)$$

Evaluation of Candidate Rules

- Combination of Measures
 - Measures are often used in combination.
 - Example: Find all rules with a minimum level of confidence then, of those rules, sort rules in descending order by lift or leverage.
- Problem: These measures do not reflect novelty of rules i.e. differentiate between known rules and rules that are new to an observer.
 - Novelty and value of rules need to be evaluated by a human observer.

Applications of Association Rules

- Market basket analysis
 - Better merchandising, Placement of products, and Promotion plan
- Recommender system
 - Discover related products or similar customers
- Clickstream analysis
 - Analyse data of web browsing and use clicks
- Much more...

An Example:

Transactions in a Grocery Store

- Employ **Apriori** algorithm to
 - **Generate** frequent itemsets and rules
 - **Visualise** the generated rules
- Use R and **arules** and **arulesViz** packages

```
install.packages('arules')  
install.packages('arulesViz')
```

```
library('arules')  
library('arulesViz')
```

An Example:

Transactions in a Grocery Store

- The **Groceries** dataset
 - 30 days of real-world sale transactions of a store

```
data(Groceries)
Groceries
transactions in sparse format with
  9835 transactions (rows) and
  169 items (columns)

summary(Groceries)
transactions as itemMatrix in sparse format with
  9835 rows (elements/itemsets/transactions) and
  169 columns (items) and a density of 0.02609146

most frequent items:
      whole milk other vegetables      rolls/buns
           2513              1903           1809
      yogurt      (Other)
           1372          34055
```

An Example:

Transactions in a Grocery Store

- Class of “**transactions**” (in **arules** package)
 - **itemsetInfo**: a **data frame** with vectors of the same length as the number of transactions
 - Say, store Customer ID
 - **itemInfo**: A **data frame** to store item labels
 - **data**: A **binary incidence matrix** that indicates which item labels appear in every transaction

```
class(Groceries)
[1] "transactions"
attr(,"package")
[1] "arules"
```

An Example:

Transactions in a Grocery Store

- **itemInfo**: A data frame to store item labels

Example: List the first 10 item labels

```
Groceries@itemInfo[1:10, "labels"]
```

```
[1] "frankfurter" "sausage" "liver loaf" "ham" "meat"  
[6] "finished products" "organic sausage" "chicken"  
"turkey" "pork"
```

An Example:

Transactions in a Grocery Store

- Example: list higher order descriptions

```
Groceries@itemInfo[1:20,]
```

	labels	level2	level1
1	frankfurter	sausage	meet and sausage
2	sausage	sausage	meet and sausage
3	liver loaf	sausage	meet and sausage
4	ham	sausage	meet and sausage
5	meat	sausage	meet and sausage
6	finished products	sausage	meet and sausage
7	organic sausage	sausage	meet and sausage
8	chicken	poultry	meet and sausage
9	turkey	poultry	meet and sausage
10	pork	pork	meet and sausage

An Example:

Transactions in a Grocery Store

- **data**: A binary incidence matrix that indicates which item labels appear in every transaction
- **Display** the 10th to 20th transactions of the dataset
- Try: `Groceries@data[1,]` (returns sparse matrix)
- Better:

```
apply(Groceries@data[,10:20], 2,  
      function(r) paste(Groceries@itemInfo[r,"labels"], collapse=", ")  
      )  
  
[1] "whole milk, cereals"  
[2] "tropical fruit, other vegetables, white bread, bottled water,  
chocolate"  
[3] "citrus fruit, tropical fruit, whole milk, butter, curd, yogurt,  
flour, bottled water, dishes"  
[4] "beef"
```

An Example:

Transactions in a Grocery Store

- Frequent Itemset Generation
 - Use `apriori()` function in the `arules` package
 - The `apriori()` function executes all the steps (What are they?) once
 - Specific the minimum support threshold
 - Until it runs out of support or until `k` (in k -itemset) reaches the default `maxlen = 10`

[illegible]

An Example:

Transactions in a Grocery Store

```
itemsets <- apriori(Groceries, parameter=list(minlen=1, support=0.02,  
                                              target="frequent itemsets"))
```

parameter specification:

confidence	minval	smax	arem	aval	originalSupport	support	minlen
0.8	0.1	1	none	FALSE	TRUE	0.02	1

maxlen	target	ext
10	frequent itemsets	FALSE

algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

```
apriori - find association rules with the apriori algorithm  
version 4.21 (2004.05.09)          (c) 1996-2004  Christian Borgelt  
set item appearances ...[0 item(s)] done [0.00s].  
set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].  
sorting and recoding items ... [59 item(s)] done [0.00s].  
creating transaction tree ... done [0.00s].  
checking subsets of size 1 2 3 done [0.00s].  
writing ... [122 set(s)] done [0.00s].  
creating S4 object  ... done [0.00s].
```


An Example:

Transactions in a Grocery Store

- **Rule Generation and Visualization**
 - Use `apriori()` function in the `arules` package

```
rules <- apriori(Groceries, parameter=list(support=0.001,  
                                           confidence=0.6, target = "rules"))
```

An Example:

Transactions in a Grocery Store

```
rules <- apriori(Groceries, parameter=list(support=0.001,  
                                           confidence=0.6, target = "rules"))
```

parameter specification:

confidence	minval	smax	arem	aval	originalSupport	support	minlen
0.6	0.1	1	none	FALSE	TRUE	0.001	1

maxlen	target	ext
10	rules	FALSE

algorithmic control:

filter	tree	heap	memopt	load	sort	verbose
0.1	TRUE	TRUE	FALSE	TRUE	2	TRUE

apriori - find association rules with the apriori algorithm

version 4.21 (2004.05.09) (c) 1996-2004 Christian Borgelt

set item appearances ...[0 item(s)] done [0.00s].

set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].

sorting and recoding items ... [157 item(s)] done [0.00s].

creating transaction tree ... done [0.00s].

checking subsets of size 1 2 3 4 5 6 done [0.01s].

writing ... [2918 rule(s)] done [0.00s].

creating S4 object ... done [0.01s].

An Example:

Transactions in a Grocery Store

```
summary(rules)
set of 2918 rules

rule length distribution (lhs + rhs):sizes
  2    3    4    5    6
  3  490 1765  626   34

  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  2.000  4.000   4.000   4.068  4.000   6.000


summary of quality measures:
      support      confidence      lift
Min.   :0.001017 Min.   :0.6000 Min.   : 2.348
1st Qu.:0.001118 1st Qu.:0.6316 1st Qu.: 2.668
Median :0.001220 Median :0.6818 Median : 3.168
Mean   :0.001480 Mean   :0.7028 Mean   : 3.450
3rd Qu.:0.001525 3rd Qu.:0.7500 3rd Qu.: 3.692
Max.   :0.009354 Max.   :1.0000 Max.   :18.996

mining info:
      data ntransactions support confidence
Groceries      9835    0.001      0.6
```

An Example:

Transactions in a Grocery Store

- Visualization: `plot(rules)` function



Scatterplot of the 2,918 rules with minimum support 0.001 and minimum confidence 0.6

An Example:

Transactions in a Grocery Store

- Visualization: `plot(rules@quality)`

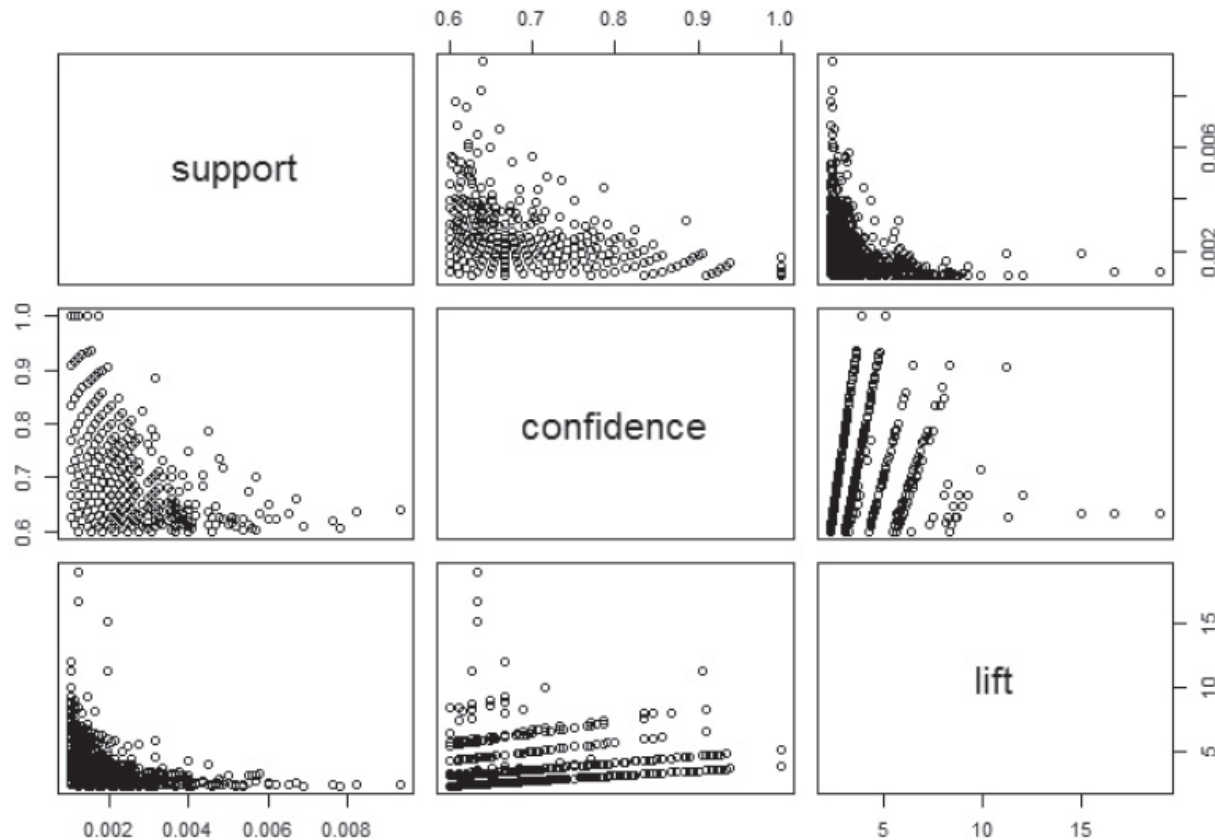



FIGURE 5-4 Scatterplot matrix on the support, confidence, and lift of the 2,918 rules

An Example: Transactions in a Grocery Store

- Visualization: `plot()`

```
confidentRules <- rules[quality(rules)$confidence > 0.9]
plot(confidentRules, method="matrix", measure=c("lift", "confidence"),
     control=list(reorder=TRUE))
```

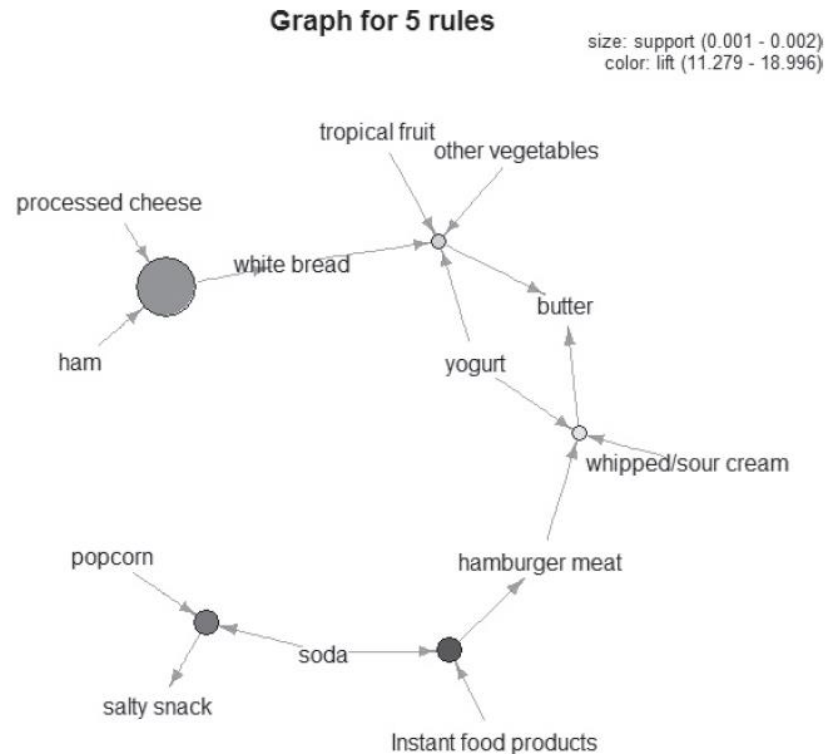


Matrix-based visualization of LHS and RHS, colored by lift and confidence

An Example: Transactions in a Grocery Store

- Visualization: `plot()`

```
highLiftRules <- head(sort(rules, by="lift"), 5)  
plot(highLiftRules, method="graph", control=list(type="items"))
```



Graph visualization of the top five rules sorted by lift

An Example:

Transactions in a Grocery Store

- Display rule content: `inspect()`

```
inspect(head(sort(rules, by="lift"), 10))
      lhs                                rhs
support confidence lift
1  {Instant food products,
    soda}                                => {hamburger meat}
0.001220132  0.6315789 18.995654
2  {soda,
    popcorn}                            => {salty snack}
0.001220132  0.6315789 16.697793
3  {ham,
    processed cheese}                  => {white bread}
0.001931876  0.6333333 15.045491
```


Validation and Testing

- Uninteresting rules
 - Involve mutually independent items
 - Cover few transactions
- Some rules could be purely coincidental
 - If 95% of customers buy X and 90% of them buy Y, then X and Y would occur together at least 85% of the time, even if there is no relationship between them
- Subjective criteria
 - Rules don't reveal unexpected profitable actions

Diagnostics

- Measures like confidence, lift, and leverage shall be used along with human insights
- Properly specify the minimum support
- Apriori algorithm can be computationally expensive!
 - Various methods to improve Apriori's efficiency

Recap: Advanced Analytical Theory and Methods: Association Rules

- Apriori Algorithm
 - Unsupervised analysis technique
 - Uncovers relationships among items
- A wide range of applications
- Several measures to help validation
- Interesting rules
 - Do not seem obvious
 - Provide valuable insights

