

Classification: Up until now

- Classification with ANNs can be computationally expensive
 - Hard to tune (number of layers, number of neurons, etc.)
 - Prone to over-fitting (lower training error but higher test error)
 - Local minimum
 - Initialisation sensitive
 - Any more?
-
- To deal with these issues we introduce

Support Vector Machine (SVM)



1: Support Vector Machines

- Background
- Linear SVM
- Soft-margin classifier
- Non-linear SVM and the kernel trick



Support Vector Machines

- Perceptron Neural Systems
 - Biologically motivated data-driven optimal (?) classifiers.
 - Massive parallel system
 - ...but are slow when implemented and run on single CPU systems.
- Support Vector Machines
 - Data-driven optimal classifiers.
 - Same computational abilities as MLP.
 - Could be much faster on single CPU systems



Terminology: “Support Vectors”

We have already learned:

- An MLP can find an optimal decision boundary for a given set of (training-)samples for which target values are available.
- Training data is available in the form of multi-dimensional fixed sized **vectors**.
- The training algorithm is entirely driven by the training data.
- Thus, one could say that the vectors in the training set **support** the MLP in finding the decision boundary.
- Note that the decision boundary can also be found if we only have had the data which are closest to the decision boundary.

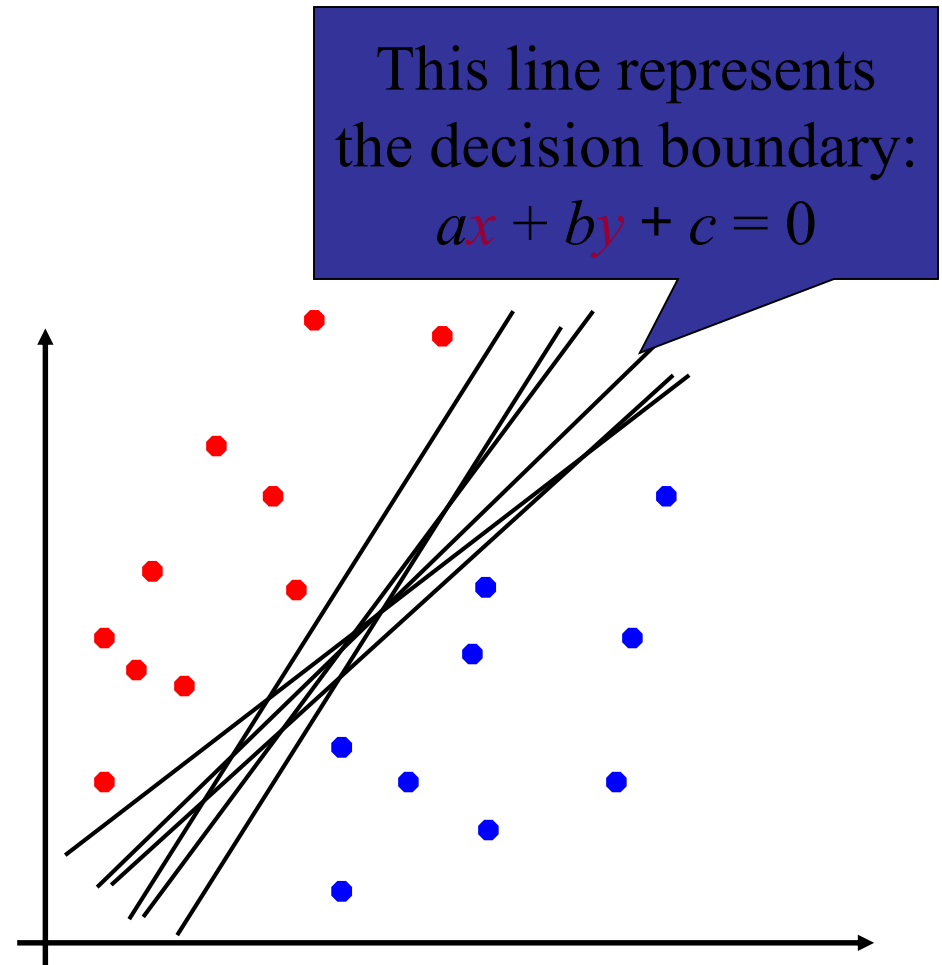
In other words:

- The **minimum sub-set of the training data** needed to find the decision boundary are called the **support vectors**.




SVM - background

- Example: 2-dimensional, binary, linearly separable classification problems.
- Decision boundary is a line $ax+by+c$.
- Task is to find a , b , and c that separates the two classes of data.
- There may be infinite many a, b, c that meet the requirement.
- **Which one is best from your point of view?**



SVM – the aim

- SVM aims at **maximizing** the distance between the decision boundary and the “difficult points” close to decision boundary




Among all the possible decisions boundaries that separate two classes, the SVM computes the one that **maximizes** the distance to the nearest data points.

Note: MLP has a different goal.

Alternative intuition

- Assume that we have “fat” separator lines. The fatter the line the less choices we have.
- SVM aims at finding the fattest possible separator line.



SVM – Formalization

For general n-dimensional learning problems we can describe the decision boundary as follows:

$$a_1x_1 + a_2x_2 + \dots + a_nx_n + b = 0$$

Or, more conveniently, in vector form as follows:

$$\mathbf{a}^T \mathbf{x} + b = 0$$

Thus, the vector **a** (and b) define a **hyperplane** (and its offset).

Formalization

- \mathbf{w} : decision hyperplane normal vector
- The corresponding unit vector is $\mathbf{w}/|\mathbf{w}|$, where $|\mathbf{w}| = \sqrt{\mathbf{w}^T \mathbf{w}}$
- \mathbf{x}_i : data point i
- y_i : class of data point i (**+1 or -1**) **NB: Not 1/0**
- r : the distance of a point from the hyperplane
- Classifier is: $f(\mathbf{x}_i) = \text{sign}(\mathbf{w}^T \mathbf{x}_i + b)$
- \mathbf{x}'_i : Point on the hyperplane nearest to \mathbf{x}_i .

Thus, \mathbf{x}' is a translation of \mathbf{x} by r :

$$\mathbf{x}' = \mathbf{x} - yr\mathbf{w}/|\mathbf{w}|$$

Geometric Margin


- This satisfies $\mathbf{w}^T \mathbf{x}' + b = 0$, and hence we can write

$$\mathbf{w}^T (\mathbf{x} - yr\mathbf{w}) + b = 0$$

- Solved for r gives: $r = y \frac{\mathbf{w}^T \mathbf{x} + b}{|\mathbf{w}|}$
- This is called the **geometric margin**.
- Note: The geometric margin is invariant to scale. We can impose any scale without affecting the geometric margin.
- It is convenient to set a scale so that the geometric margin of any data point is at least 1.

Geometric Margin

- Examples closest to the hyperplane are *support vectors*.
- *Margin* ρ of the separator is the width of separation between support vectors of classes.



Derivation of finding r :

Dotted line $\mathbf{x}' - \mathbf{x}$ is perpendicular to decision boundary so parallel to \mathbf{w} .

Unit vector is $\mathbf{w}/|\mathbf{w}|$, so line is $r\mathbf{w}/|\mathbf{w}|$.

$\mathbf{x}' = \mathbf{x} - yr\mathbf{w}/|\mathbf{w}|$.

\mathbf{x}' satisfies $\mathbf{w}^T \mathbf{x}' + b = 0$.

So $\mathbf{w}^T(\mathbf{x} - yr\mathbf{w}/|\mathbf{w}|) + b = 0$

Recall that $|\mathbf{w}| = \sqrt{\mathbf{w}^T \mathbf{w}}$.

So $\mathbf{w}^T \mathbf{x} - yr|\mathbf{w}| + b = 0$

So, solving for r gives:

$$r = y(\mathbf{w}^T \mathbf{x} + b)/|\mathbf{w}|$$

Linear SVM Mathematically

The linearly separable case

- Assume that all data is at least distance 1 from the hyperplane, then the following two constraints follow for a training set $\{(\mathbf{x}_i, y_i)\}$

$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 \quad \text{if } y_i = 1$$

$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 \quad \text{if } y_i = -1$$

- For support vectors, the inequality becomes an equality
- Then, since each example's distance from the hyperplane is

$$r = y \frac{\mathbf{w}^T \mathbf{x} + b}{|\mathbf{w}|}$$

the margin is: $\rho = \frac{2}{|\mathbf{w}|}$




Linear Support Vector Machine (SVM)

- **Hyperplane**

$$\mathbf{w}^T \mathbf{x} + b = 0$$

- **Extra scale constraint:**

$$\min_{i=1,\dots,n} |\mathbf{w}^T \mathbf{x}_i + b| = 1$$



Linear SVMs Mathematically

- Then we can formulate the *quadratic optimization problem*:

Find \mathbf{w} and b such that

$$\rho = \frac{2}{|\mathbf{w}|} \text{ is maximized; and for all } \{(\mathbf{x}_i, y_i)\}$$

$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 \text{ if } y_i = 1; \quad \mathbf{w}^T \mathbf{x}_i + b \leq -1 \text{ if } y_i = -1$$

- A better formulation ($\min |\mathbf{w}| = \max 1/|\mathbf{w}|$):

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \text{ is minimized;}$$

$$\text{and for all } \{(\mathbf{x}_i, y_i)\}: \quad y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$



Solving the Optimization Problem

Find \mathbf{w} and b such that
 $\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$ is minimized;
and for all $\{(\mathbf{x}_i, y_i)\}$: $y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$

- This is now optimizing a **quadratic function subject to linear constraints**
- Quadratic optimization problems are a well-known class of mathematical programming problem, and many (intricate) algorithms exist for solving them (with many special ones built for SVMs)
- The solution involves constructing a *dual problem* where a *Lagrange multiplier* α_i is associated with every constraint in the primary problem.
- These methods are described in subjects on advanced math. The particulars will be omitted here. Instead we refer to **readily available software** that can solve quadratic programming problems very efficiently. See [LIBSVM](#).



Solving the Optimization Problem

- Find the optimal \mathbf{w} and b

$$\{\mathbf{w}^*, b^*\} = \min \frac{1}{2} \|\mathbf{w}\|^2$$

Subject to $y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, i = 1, 2, \dots, n$

- The primal problem
 - The cost function $\|\mathbf{w}\|^2/2$ is a convex function
 - The constraints are linear in \mathbf{w}

- Lagrangian function

$$J(\mathbf{w}, b, \alpha) = \frac{1}{2} \mathbf{w}^\top \mathbf{w} - \sum_{i=1}^n \alpha_i [y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1]$$

- Two conditions of optimality

$$\frac{\partial J(\mathbf{w}, b, \alpha)}{\partial \mathbf{w}} = \mathbf{0} \implies \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i; \quad \frac{\partial J(\mathbf{w}, b, \alpha)}{\partial b} = \mathbf{0} \implies \sum_{i=1}^n \alpha_i y_i = 0$$



Solving the Optimization Problem

- Kuhn-Tucker conditions

$$\alpha_i [y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1] = 0, \quad i = 1, 2, \dots, n$$

- Only the training samples satisfy $y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1 = 0$ can have nonzero α , and they are called “support vectors”
- The dual problem (find the optimal α , a QP problem)

$$\{\alpha^*\} = \max \left[\sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i^\top \mathbf{x}_j \right]$$

Subject to 1) $\sum_{i=1}^n \alpha_i y_i = 0$

2) $\alpha_i \geq 0, \quad i = 1, 2, \dots, n$

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i^* y_i \mathbf{x}_i \quad b^* = 1 - \mathbf{w}^* \mathbf{x}^s$$


$$g(\mathbf{x}) = \mathbf{w}^{*\top} \mathbf{x} + b^*$$

A support vector



Soft Margin Classification

- If the training data is **not linearly separable**, *slack variables* ξ_i can be added to **allow misclassification of difficult or noisy examples**.
- Allow some errors: Let some points be moved to where they belong, at a cost
- Still, try to minimize training set errors, and to place hyperplane “far” from each class (large margin)



Soft Margin Classification Mathematically

- The **old formulation**:

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \text{ is minimized and for all } \{(\mathbf{x}_i, y_i)\}$$
$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

- The **new formulation** incorporating slack variables:

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum \xi_i \text{ is minimized and for all } \{(\mathbf{x}_i, y_i)\}$$
$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \text{ and } \xi_i \geq 0 \text{ for all } i$$

- Parameter C can be viewed as a way to control overfitting
 - A regularization term




Classification with SVMs

- Given a new point \mathbf{x} , we can score its projection onto the hyperplane normal:
 - I.e., compute score: $\mathbf{w}^T \mathbf{x} + b$
 - **Decide class based on whether “<” or “>” 0**
 - Can set confidence threshold t .

Score $> t$: yes

Score $< -t$: no

else: don't know




Linear SVMs: Summary


- The classifier is a *separating **hyperplane***.
- The most “important” training points are the **support vectors**; they define the hyperplane.
- Quadratic optimization algorithms can identify which training points \mathbf{x}_i are support vectors.

Non-linear SVMs


- Datasets that are linearly separable (with some noise) work out great:





- But what are we going to do if the dataset is just too hard?



- How about ... **mapping data to a higher-dimensional space:**








Non-linear SVMs: Kernel trick

- General idea: The original feature space can always be mapped to some higher-dimensional feature space where the training set is separable:



The “Kernel Trick”

- The linear classifier relies on an inner product between vectors $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
- If every datapoint is mapped into high-dimensional space via some transformation $\Phi: \mathbf{x} \rightarrow \phi(\mathbf{x})$, the inner product becomes:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

- A *kernel function* is some function that corresponds to an inner product in some expanded feature space.



The “Kernel Trick”

- **Example:**

2-dimensional vectors $\mathbf{x}=[x_1 \ x_2]$; let $K(\mathbf{x}_i, \mathbf{x}_j)=(1 + \mathbf{x}_i^T \mathbf{x}_j)^2$,

Need to show that $K(\mathbf{x}_i, \mathbf{x}_j)=\varphi(\mathbf{x}_i)^T \varphi(\mathbf{x}_j)$:

$$\begin{aligned} K(\mathbf{x}_i, \mathbf{x}_j) &= (1 + \mathbf{x}_i^T \mathbf{x}_j)^2 = 1 + x_{i1}^2 x_{j1}^2 + 2 x_{i1} x_{j1} x_{i2} x_{j2} + x_{i2}^2 x_{j2}^2 + \\ & 2 x_{i1} x_{j1} + 2 x_{i2} x_{j2} = \\ & = [1 \ x_{i1}^2 \ \sqrt{2} x_{i1} x_{i2} \ x_{i2}^2 \ \sqrt{2} x_{i1} \ \sqrt{2} x_{i2}]^T [1 \ x_{j1}^2 \ \sqrt{2} x_{j1} x_{j2} \\ & x_{j2}^2 \ \sqrt{2} x_{j1} \ \sqrt{2} x_{j2}] \\ & = \varphi(\mathbf{x}_i)^T \varphi(\mathbf{x}_j) \quad \text{where } \varphi(\mathbf{x}) = [1 \ x_1^2 \ \sqrt{2} x_1 x_2 \ x_2^2 \ \sqrt{2} x_1 \\ & \sqrt{2} x_2] \end{aligned}$$




Kernels

- Why use kernels?
 - Make non-separable problem separable (in another space).
 - Map data into better representational space
- Common kernels
 - Linear kernel $\mathbf{K}(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z}$
 - Polynomial kernel $\mathbf{K}(\mathbf{x}, \mathbf{z}) = (1 + \mathbf{x}^T \mathbf{z})^d$
 - Gives feature conjunctions
 - Radial basis function (infinite dimensional space)

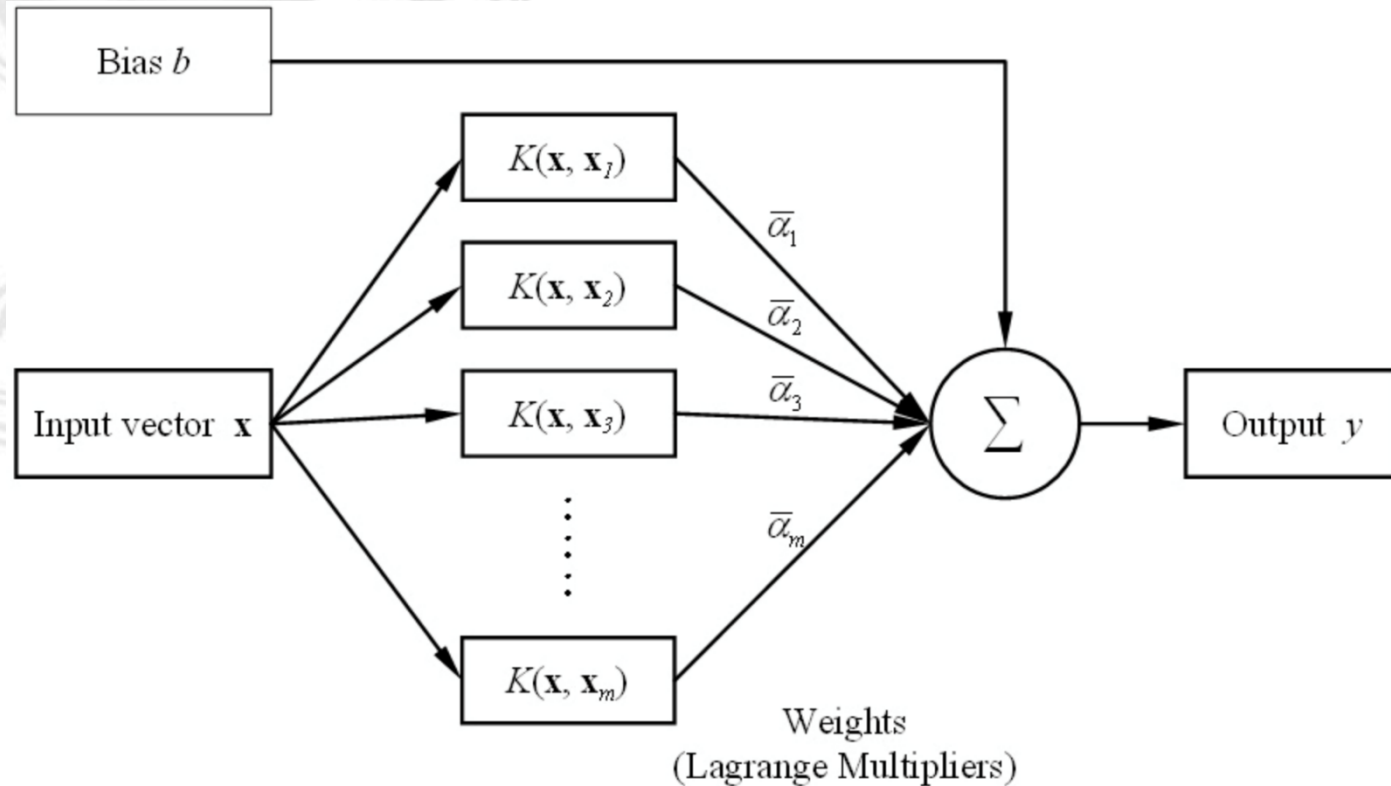
$$K(\mathbf{x}_i, \mathbf{x}_j) = e^{-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma^2}$$

Summary

- SVMs maximize the *margin* around the separating hyperplane.
 - A.k.a. large margin classifiers
- The decision function is fully specified by a subset of training samples, the *support vectors*.
- Solving SVMs is a *quadratic programming* problem
- Popularly applied on single CPU systems.



Summary



Hidden Nodes
(Support Vectors)

$$y = f(\mathbf{x}) = \sum_{k=1}^m \bar{\alpha}_k \cdot K(\mathbf{x}, \mathbf{x}_k) + b$$